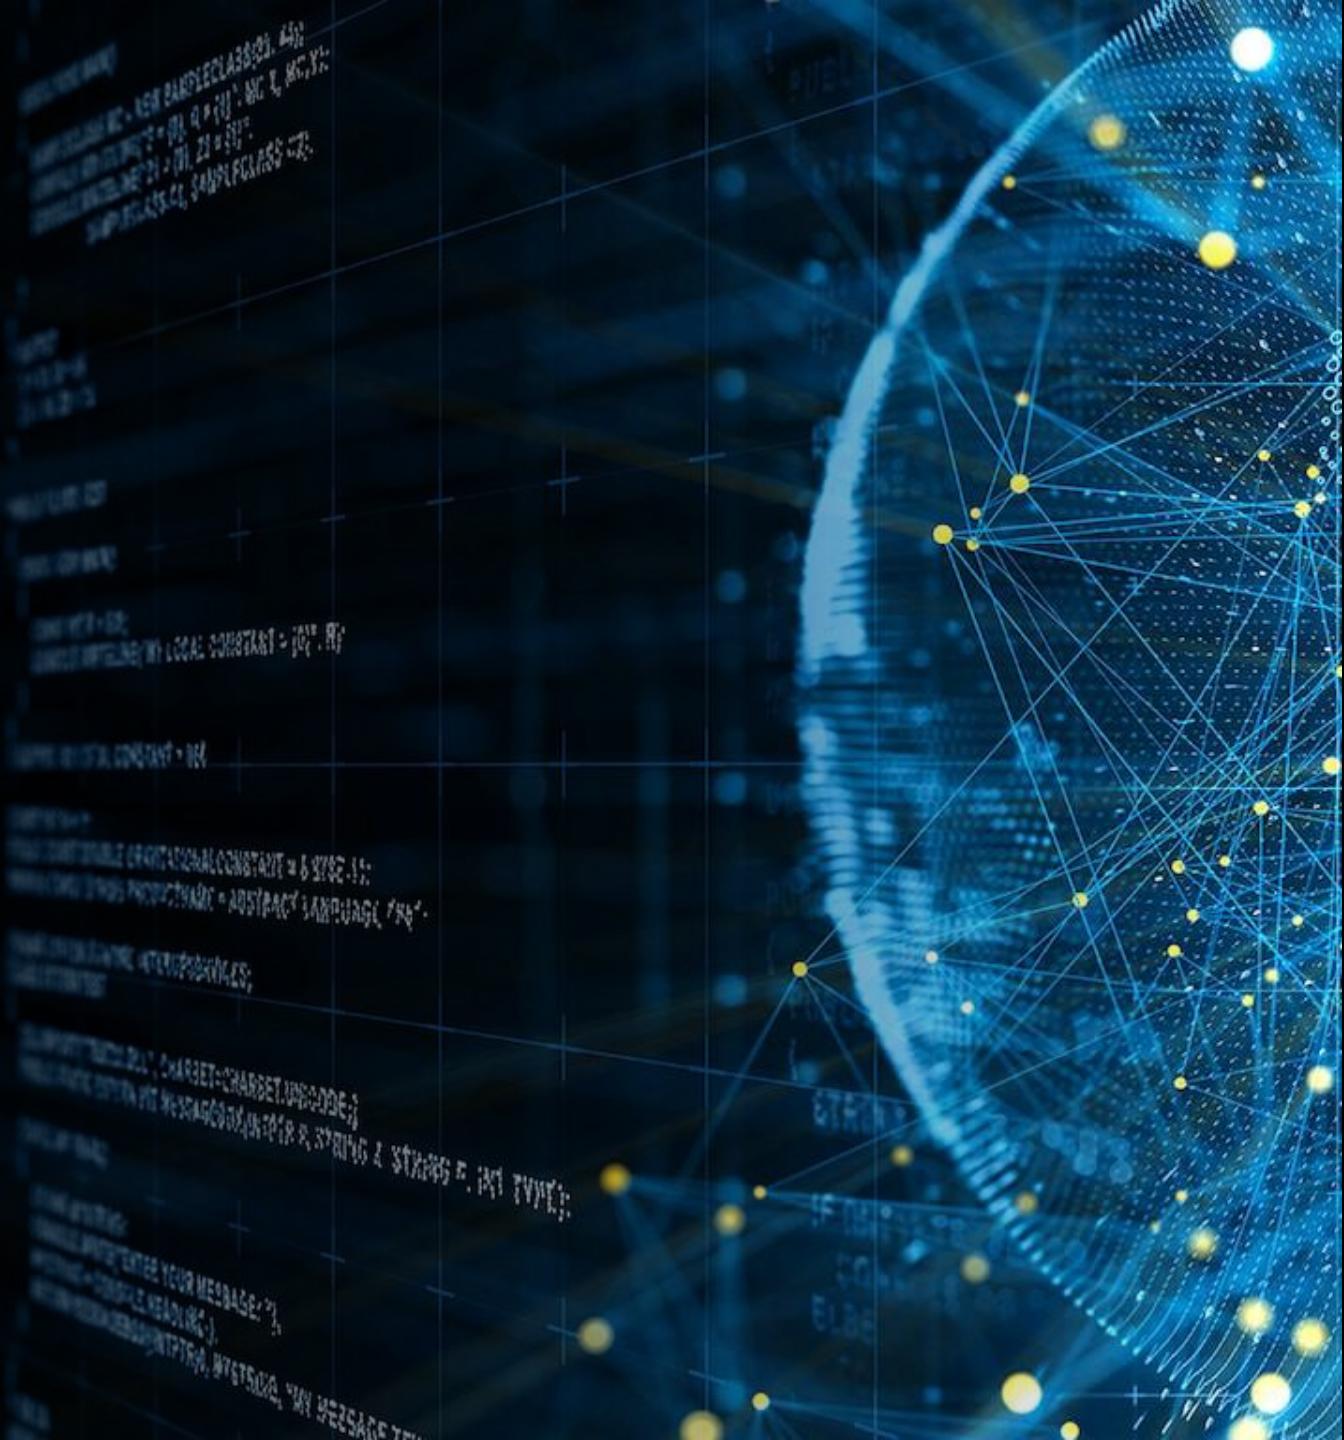


WhatTheHack

Azure OpenAI Apps



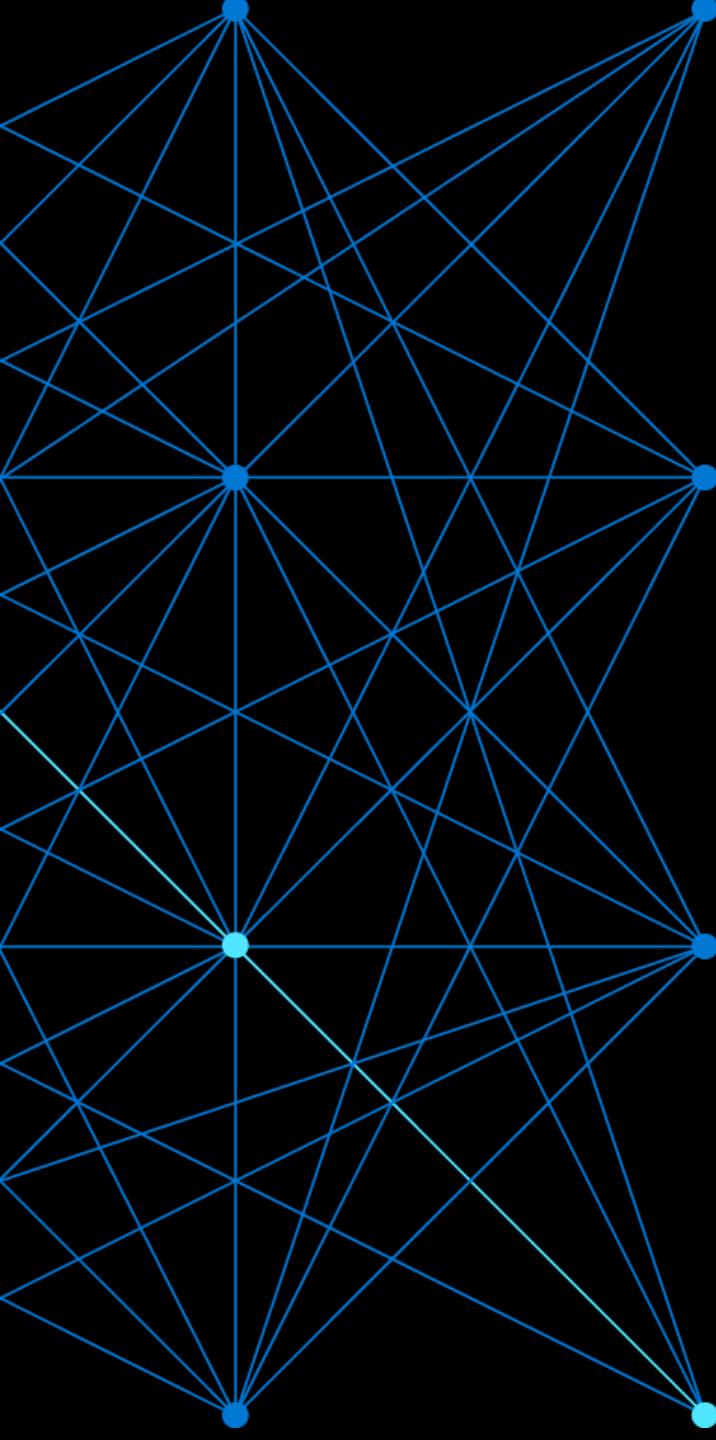
Agenda

01	Introduction
02	Deploy Azure Prerequisite Infrastructure (30-45 minutes)
03	Challenge #1 (1 hour)
04	Challenge #2 (1 hour)
05	Lunch (1 hour)
06	Challenge #3 (1-2 hours)
07	Go Home
08	Come Back
09	Challenge #4 (1-2 hours)
10	Challenge #5 (1-2 hours)

WhatTheHack Format

- Self-guided
- Self-paced (don't wait for me!)
- Group up with neighbors, coworkers
- Answers are not provided!*
- Explore, get hands on, ask questions, have fun

* but answer key is available

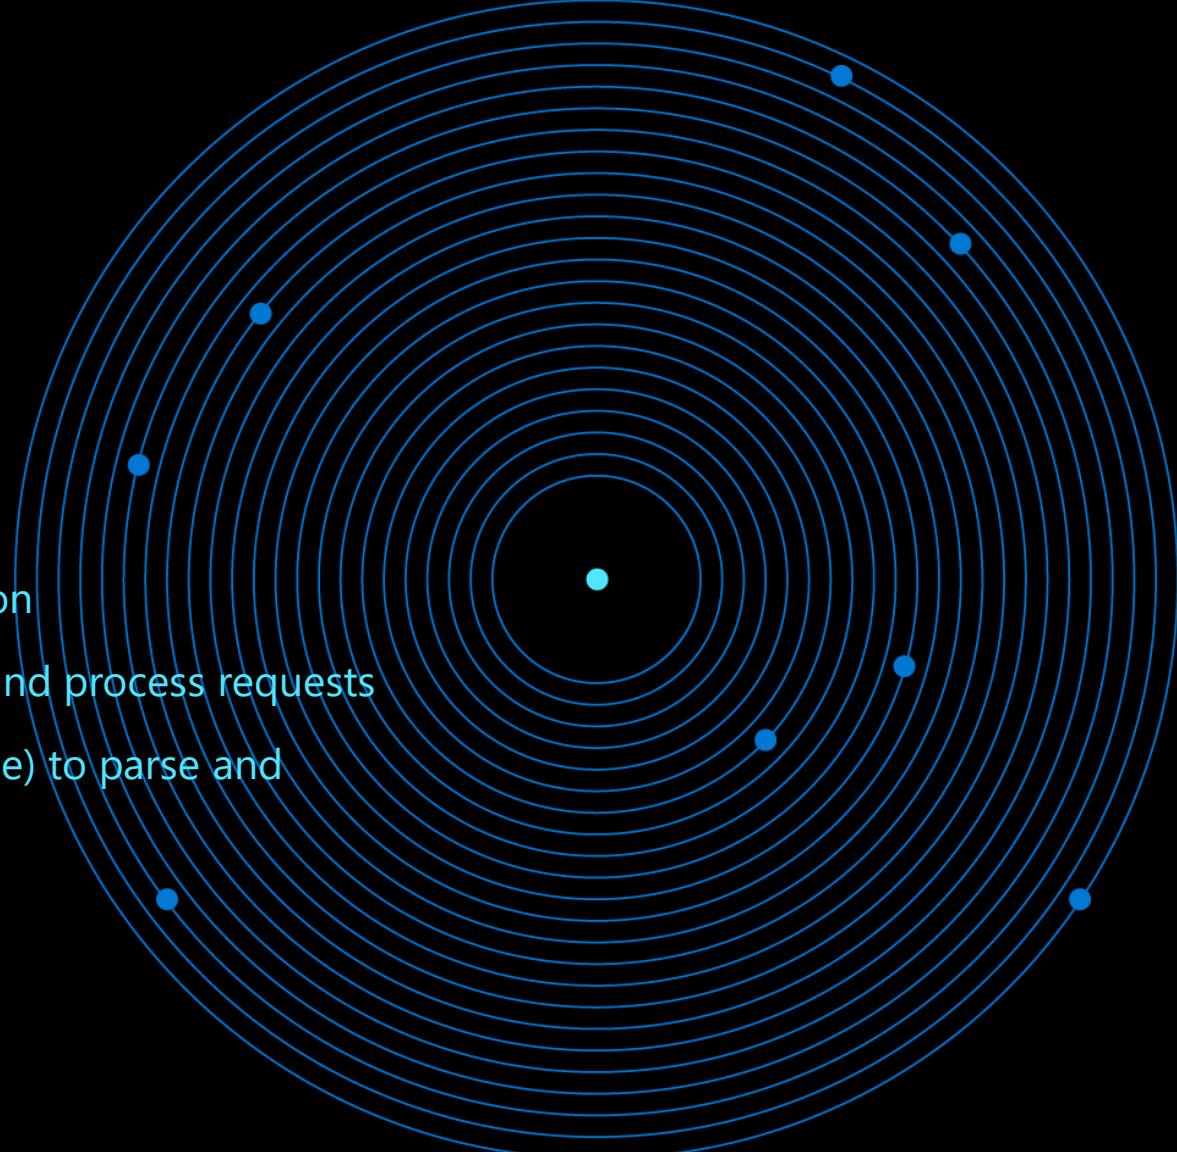


Learning Objectives

- Choose the right architecture and Azure services for different AI scenarios
- Provision, configure, and monitor Azure OpenAI resources
- Build intelligent apps and Q&A assistants using RAG architectures
- Implement solutions for both batch and real-time AI use cases

Tasks

- Deploy Azure services using infrastructure as code
- Load source data into a retrieval system through automation
- Prompt engineering to enable agents to invoke functions and process requests
- Configure Azure AI platform services(Document Intelligence) to parse and prepare the data for the LLM
- Optimization of Azure OpenAI LLMs through caching
- Understand design concerns in a generative AI application



GitHub Codespaces

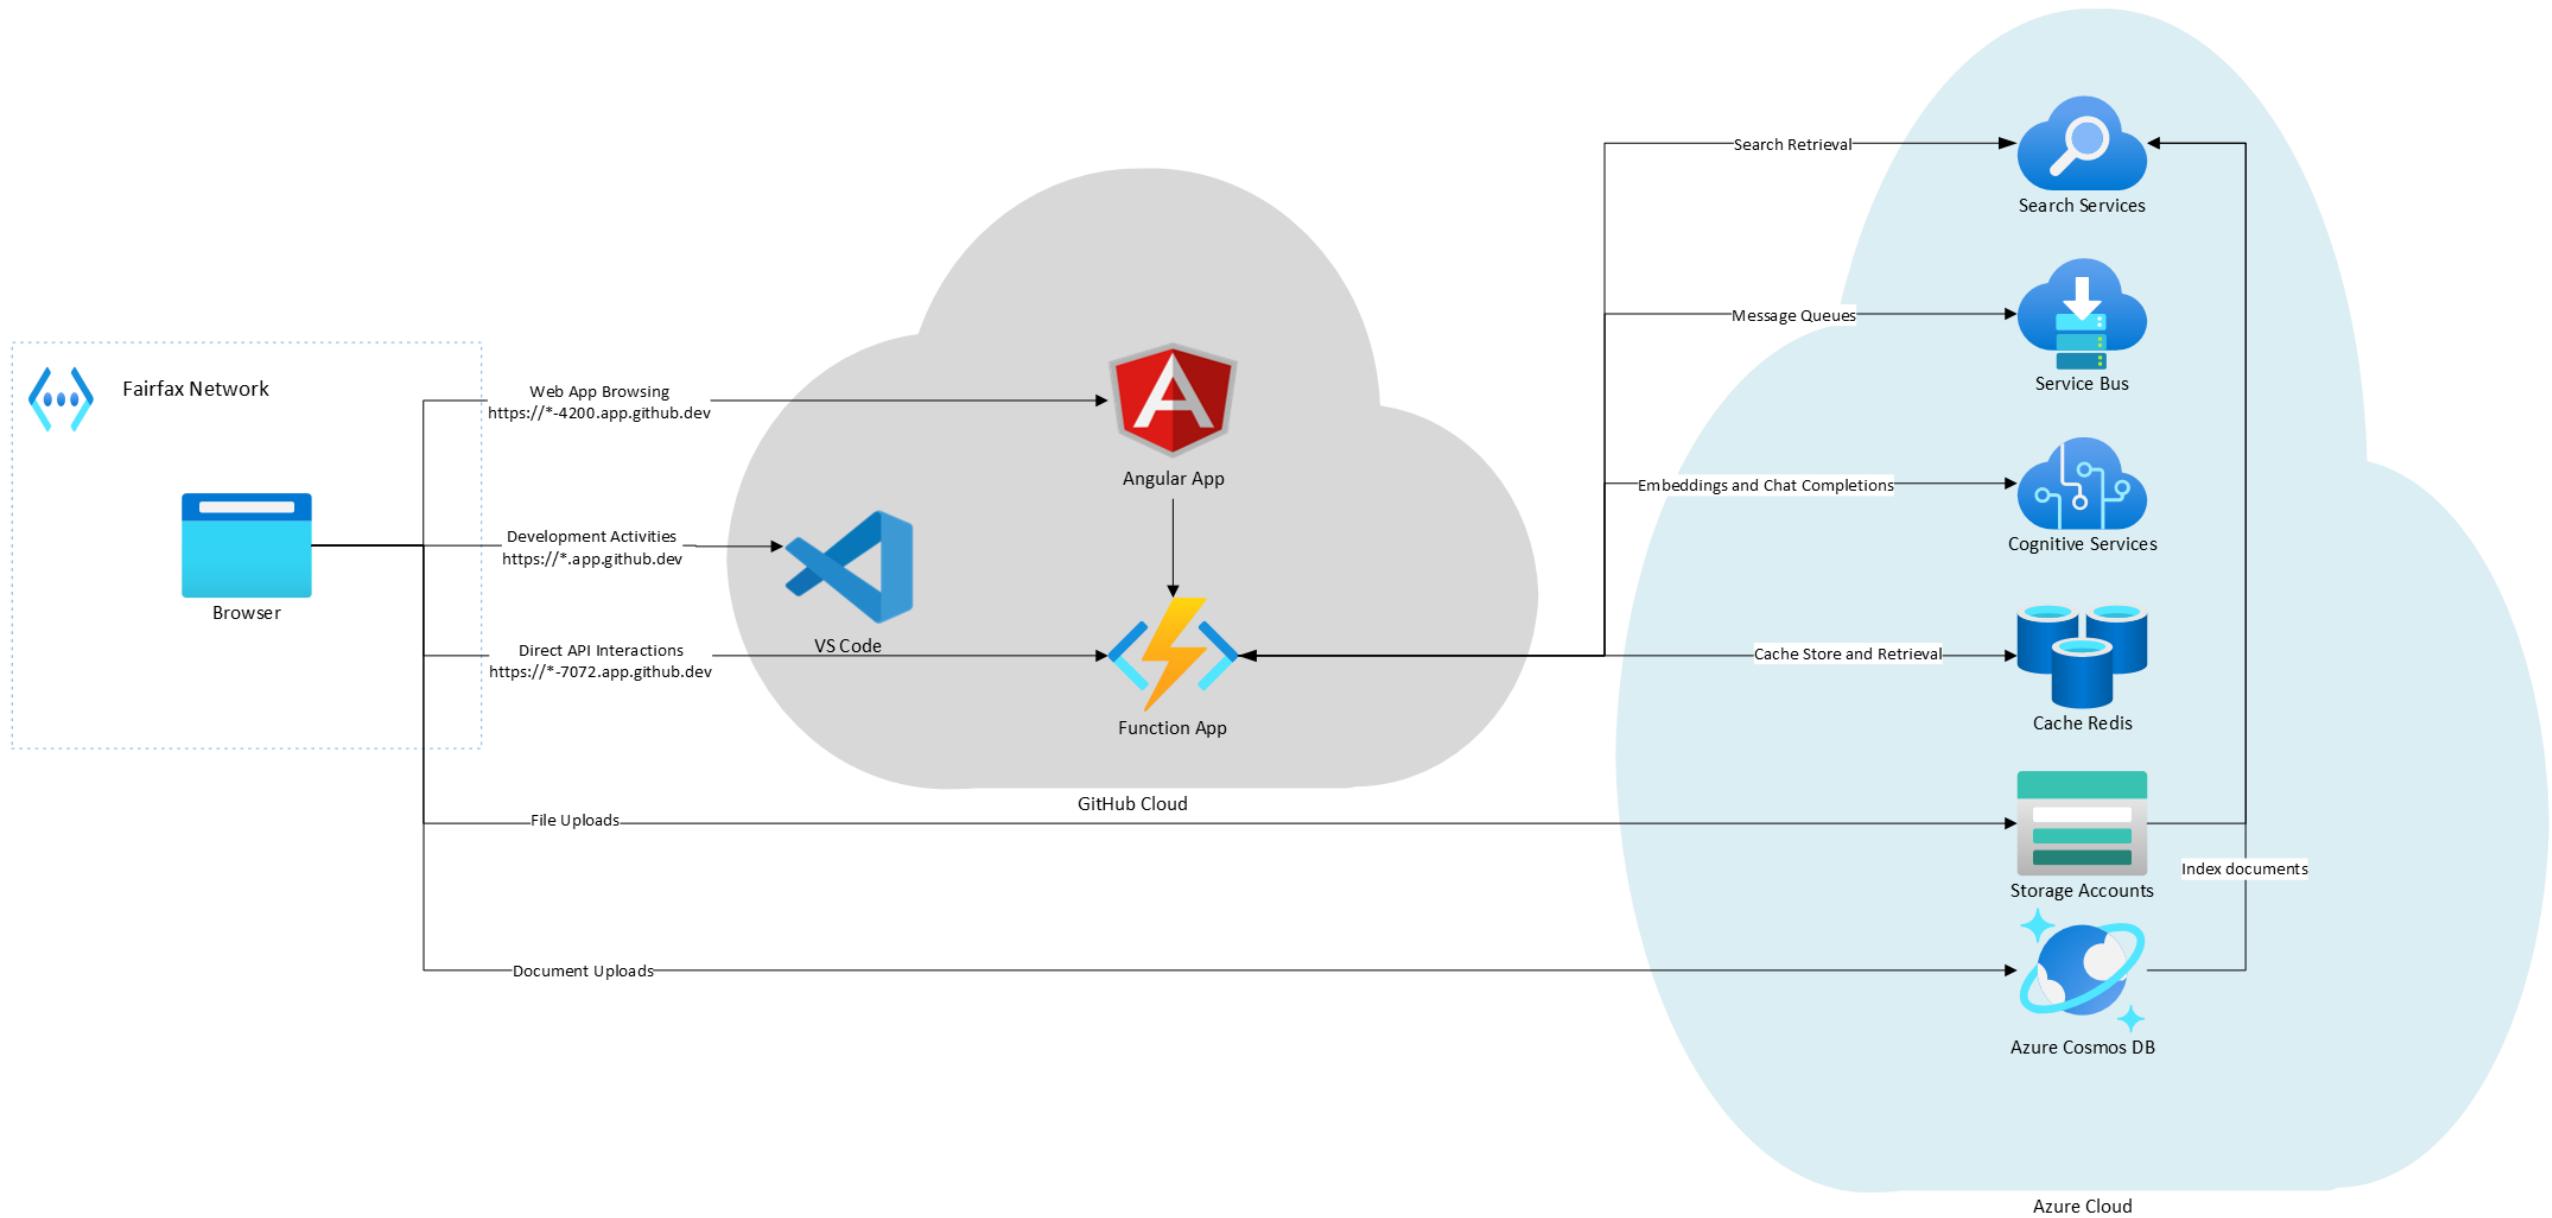
- Browser-based VS Code IDE
- Execution, runtime happens on cloud compute
- Packages all development tools, prerequisites
- Free tier for personal accounts up to 80 core hours

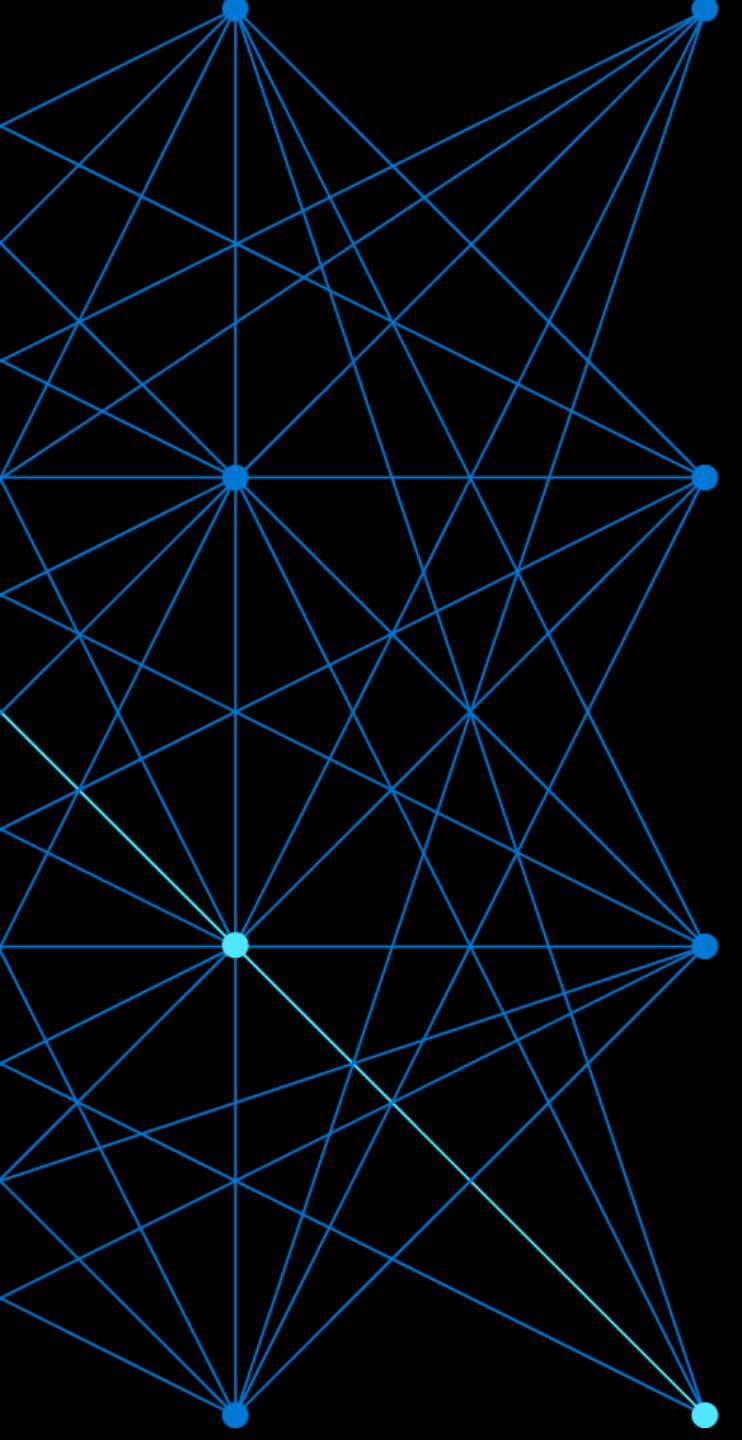


```
mona-github-github-g59jpq2w5w7.github.dev
```

The screenshot shows the GitHub Codespaces interface. At the top, there's a navigation bar with back/forward arrows, a lock icon, and a search bar. Below it is the title "mona-github-github-g59jpq2w5w7.github.dev". The main area is divided into several sections:

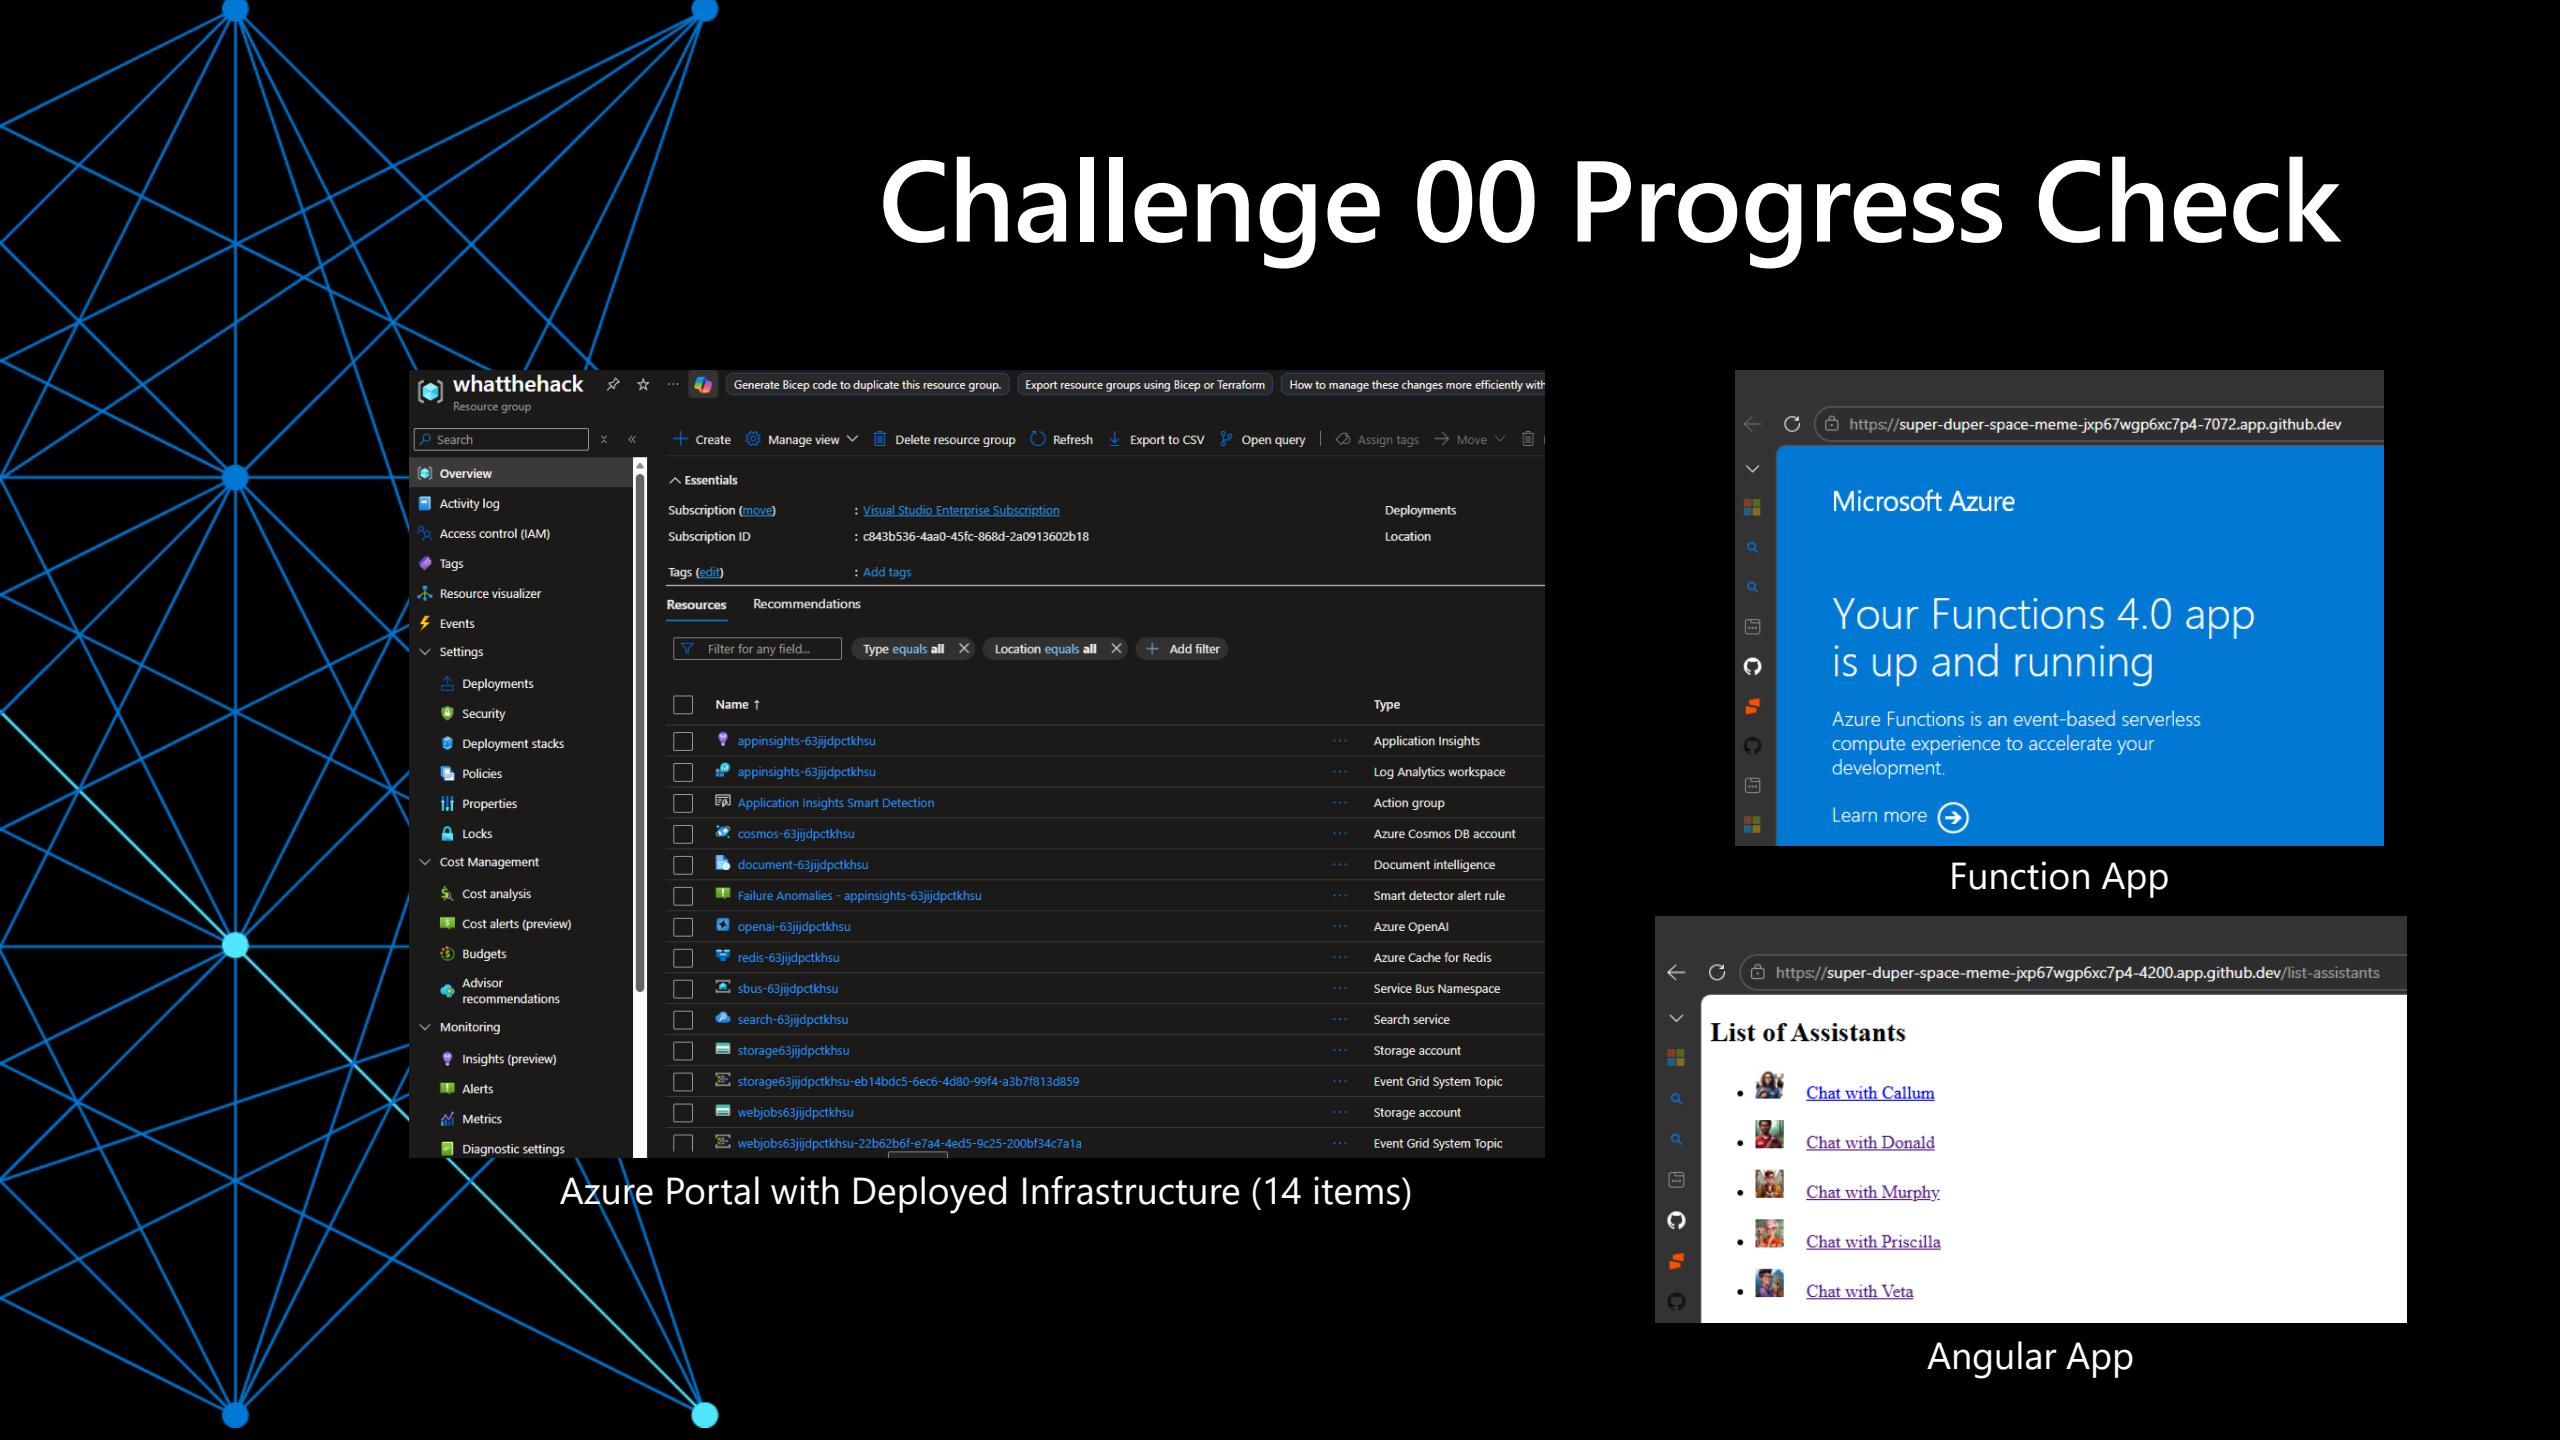
- EXPLORER:** Shows a tree view of the repository structure under "GITHUB". The "emotion.rb" file is currently selected.
- CODE EDITOR:** Displays the content of "emotion.rb". The code includes methods for handling emoji reactions and initializing content with labels and characters.
- TERMINAL:** Shows command-line logs from a session named "zsh". It lists various language server processes starting and finishing.
- STATUS BAR:** At the bottom, it shows "Ln 29, Col 7" and other terminal settings like "Tab Size: 4", "UTF-8", "LF", "Ruby", and icons for clipboard and file operations.

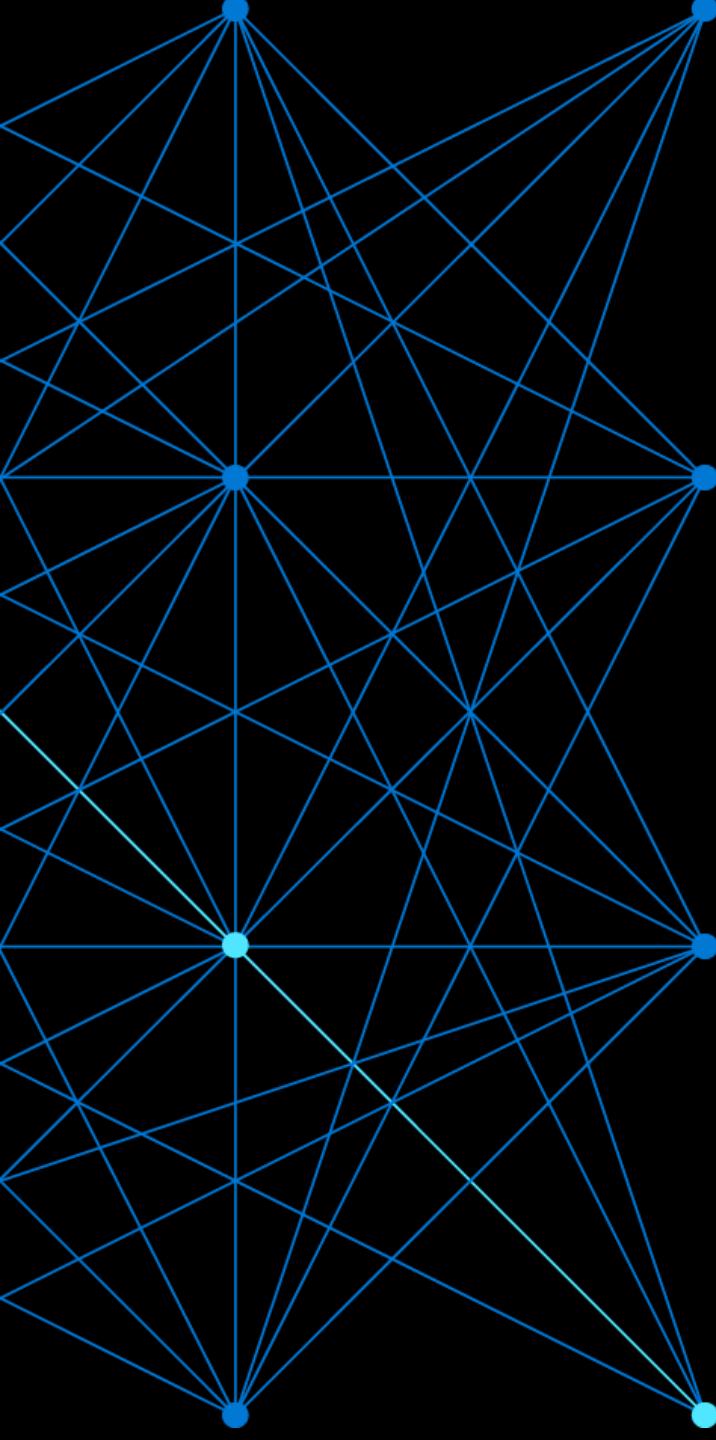




Challenge 00

Challenge 00 Progress Check



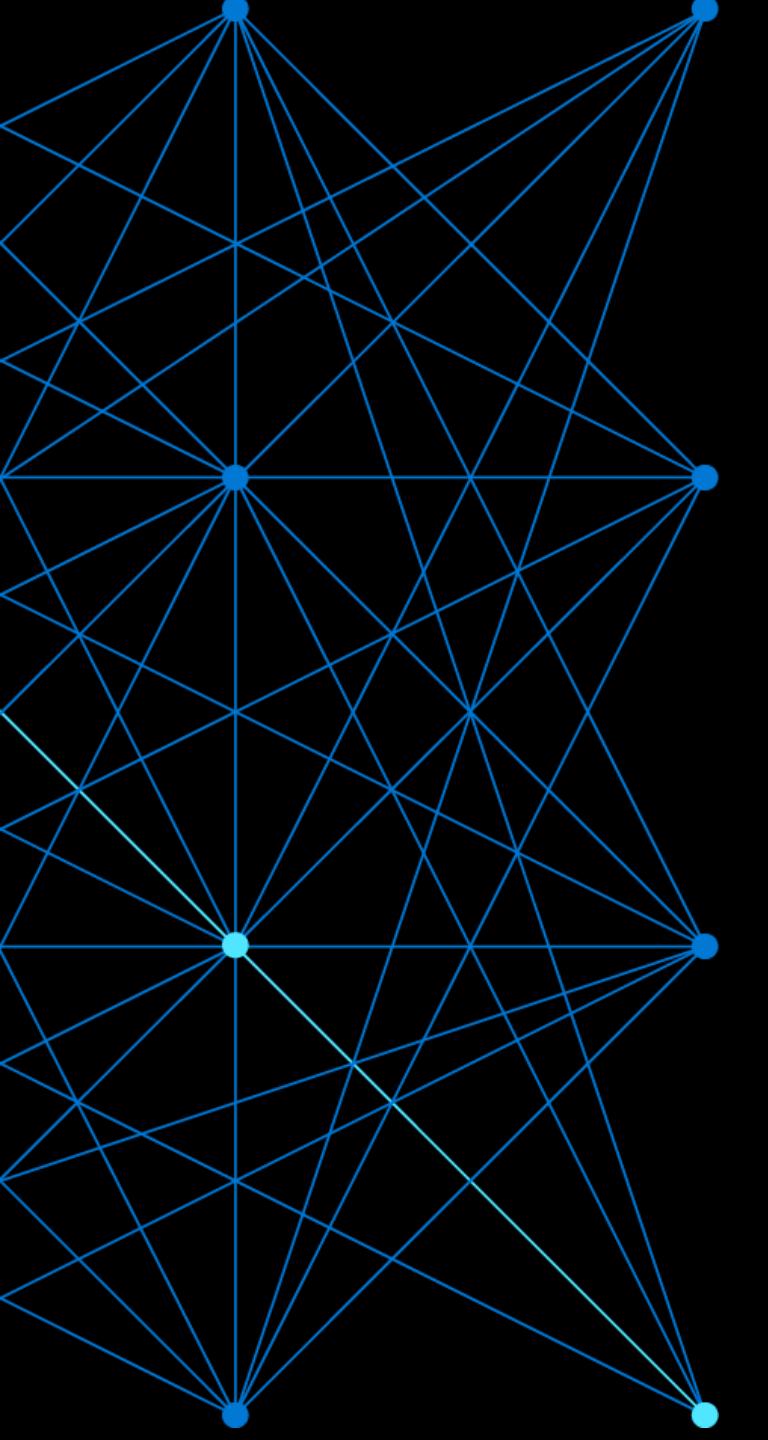


Infrastructure as Code



- Declarative language
- Azure-specific (Terraform, AWS CloudFormation, Pulumi)
- Benefits
 - Repeatable
 - Disaster recovery
 - Cost savings
 - Documentation

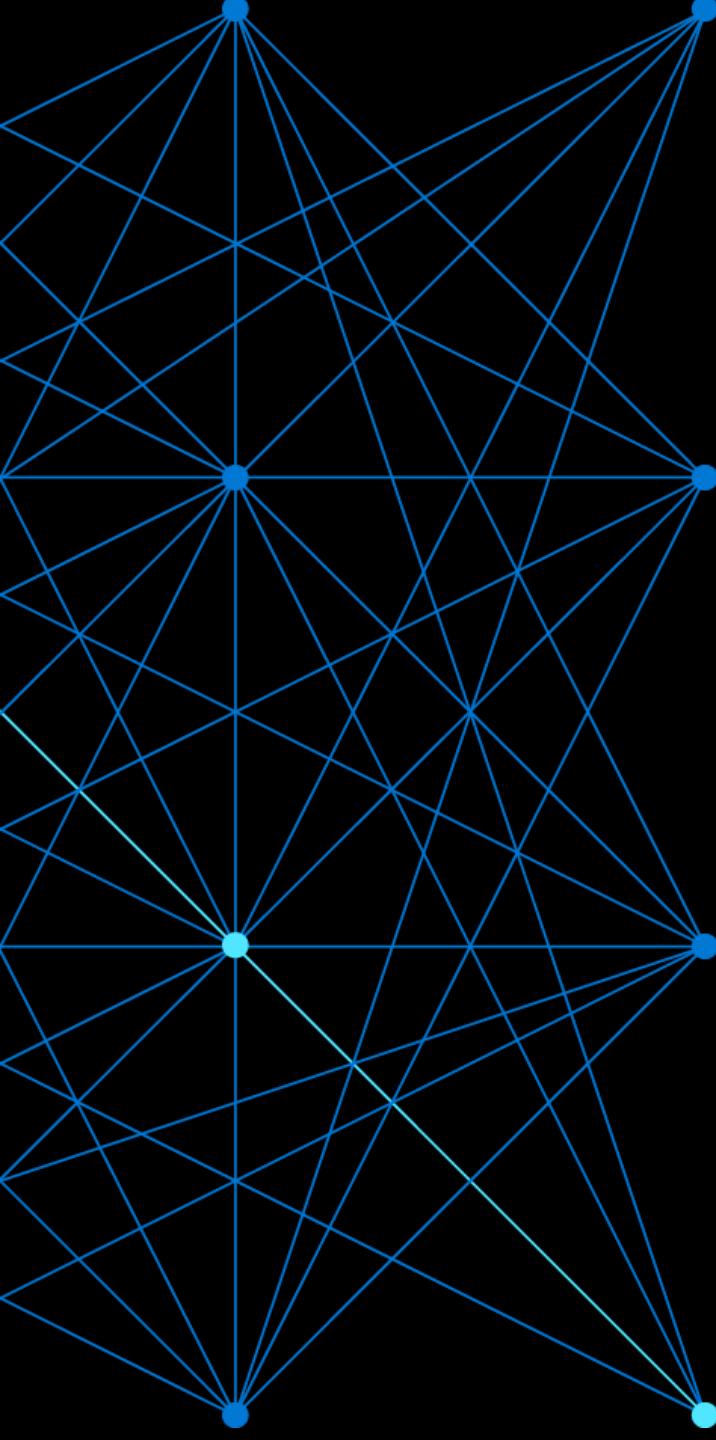
```
20      output redisPrimaryKey string = redis.outputs.primaryKey
21      output redisHostname string = redis.outputs.hostname
22
23
24      module serviceBus 'modules/servicebus.bicep' = {
25          name: 'serviceBusDeployment'
26          params: {
27              name: 'sbus-${suffix}'
28              location: location
29              queues: ['orange', 'lemon', 'grapefruit', 'tangerine']
30          }
31      }
32
33      output serviceBusConnectionString string = serviceBus.outputs.connectionString
34
35      module cosmos 'modules/cosmos.bicep' = {
36          name: 'cosmosDeployment'
37          params: {
38              name: 'cosmos-${suffix}'
39              databaseName: 'contoso'
40              location: location
41              containers: [
42                  { name: 'yachts',           partitionKey: '/yachtId' },
43                  { name: 'customers',        partitionKey: '/email' },
44                  { name: 'reservations',    partitionKey: '/id' },
45                  { name: 'examsubmissions', partitionKey: '/id' }
46              ]
47          }
48      }
```



Challenge 01

Discussion Points

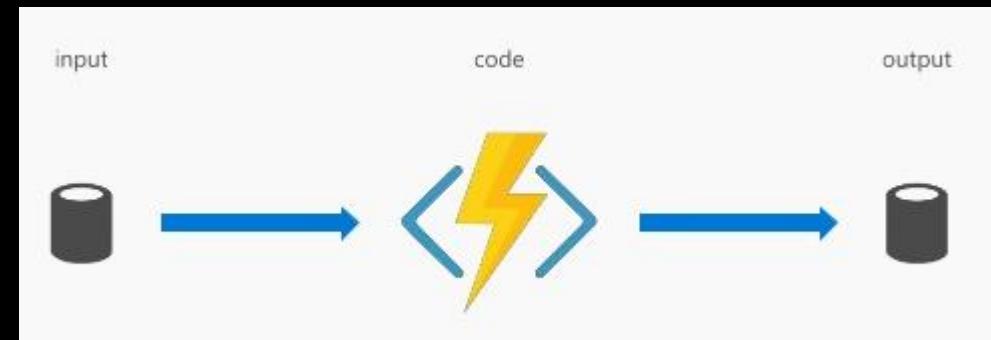
- Azure CLI
- REST Client
- Function Bindings
- Azure CosmosDB, Storage, AI Search

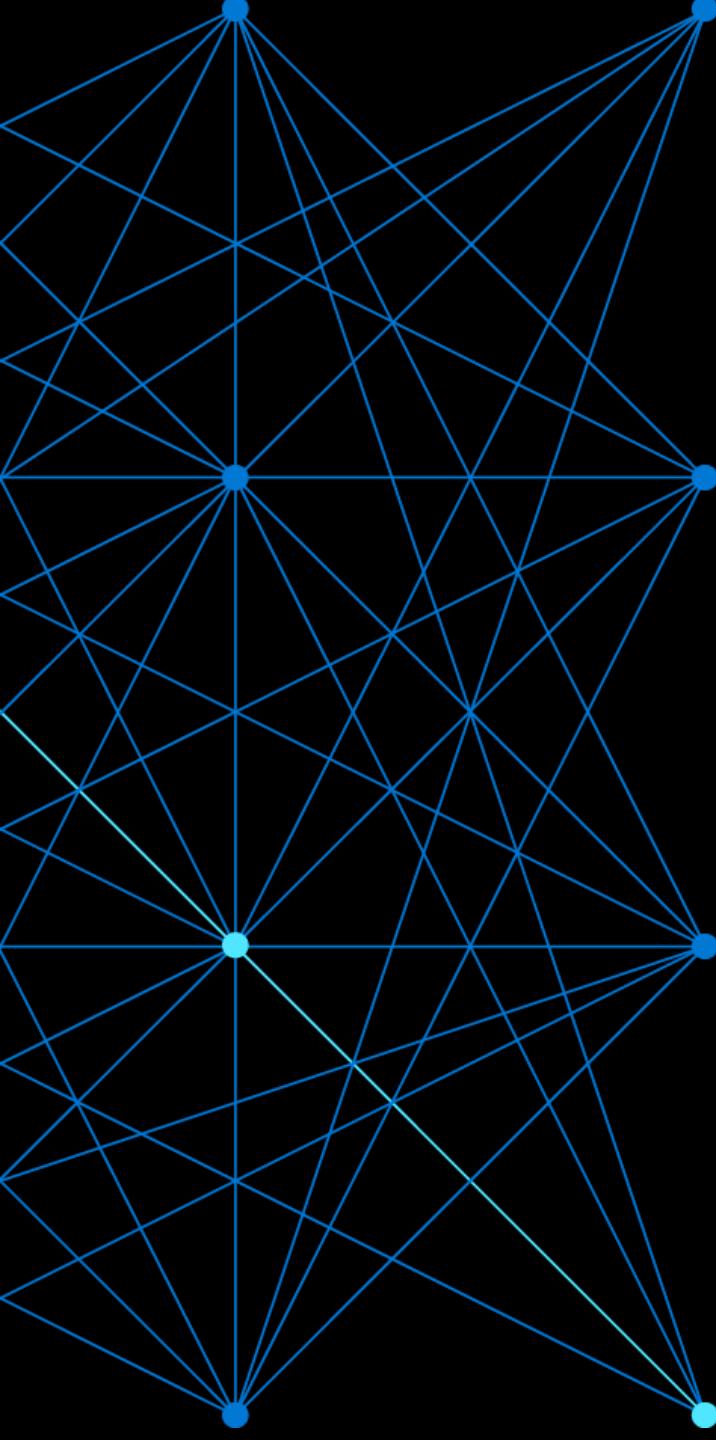


Azure Functions



- Functions-as-a-Service
- .NET, Python, Javascript, Java
- Local emulator available
- Pay-per-execution or hosted plans
- Bindings for
 - HTTP
 - Queues
 - Databases
 - Timers
 - and more

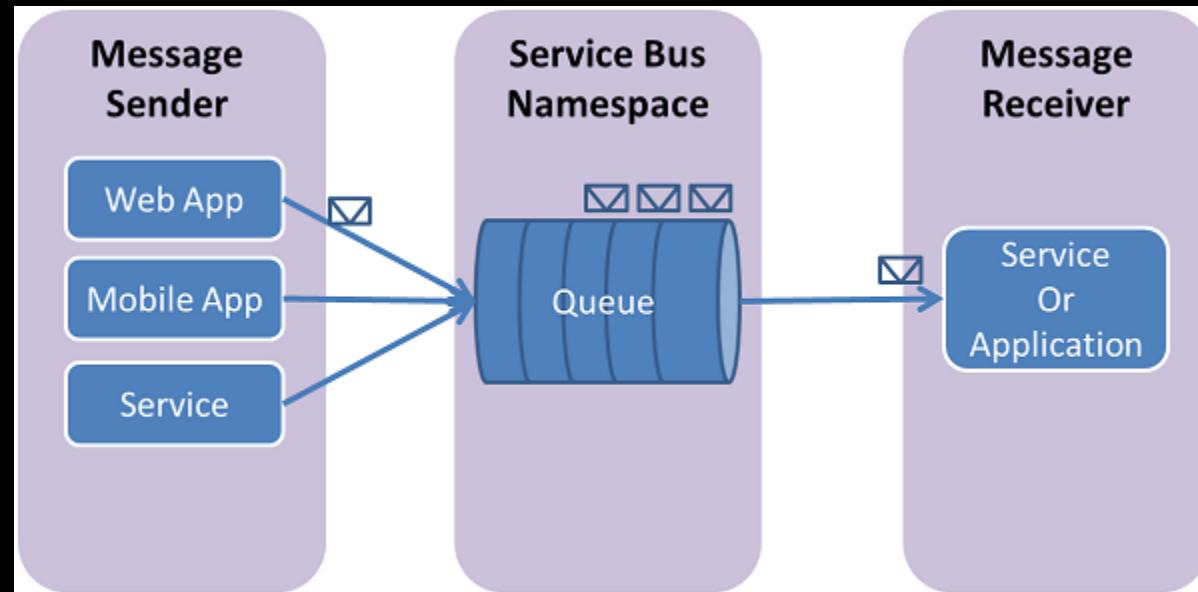


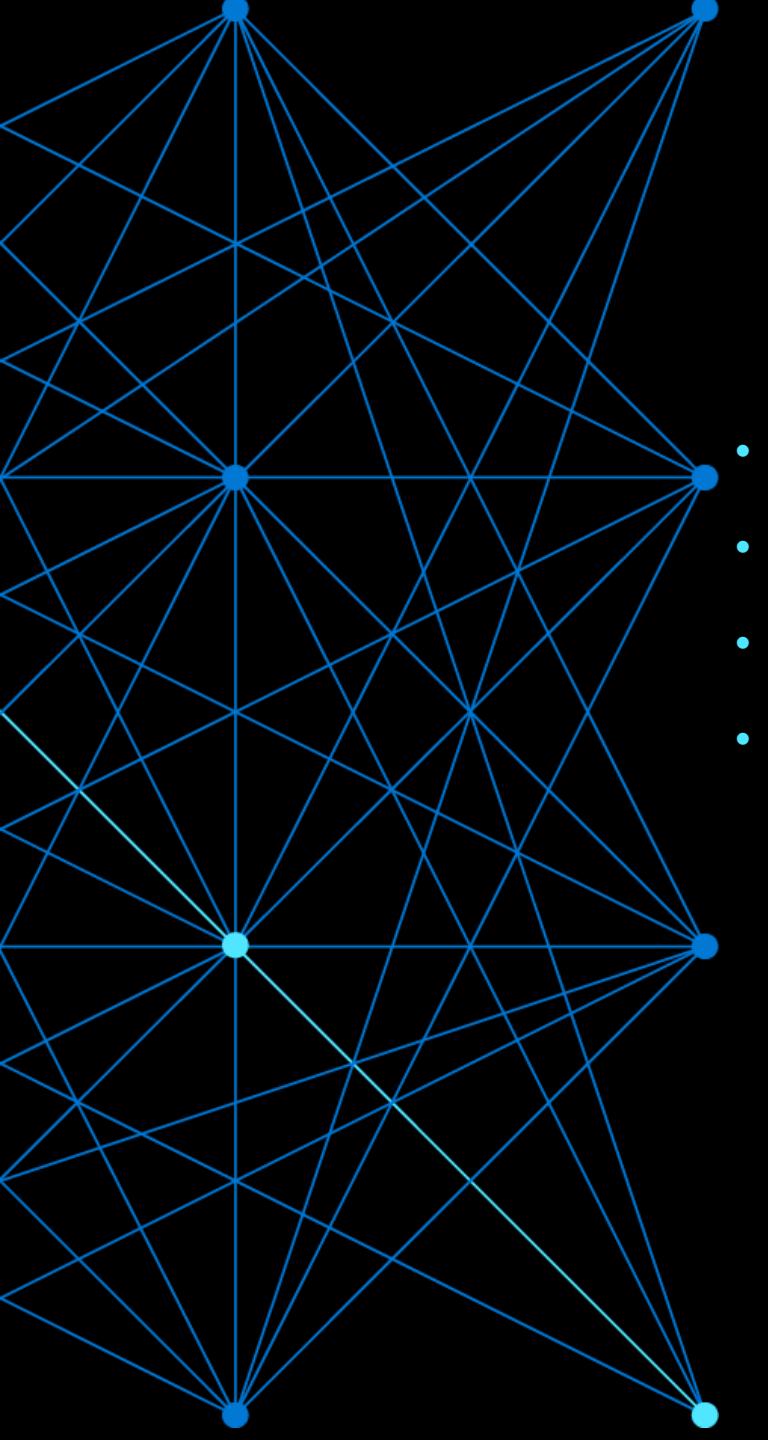


Azure Service Bus



- Platform-as-a-Service
- Queues
- Topics and Subscriptions

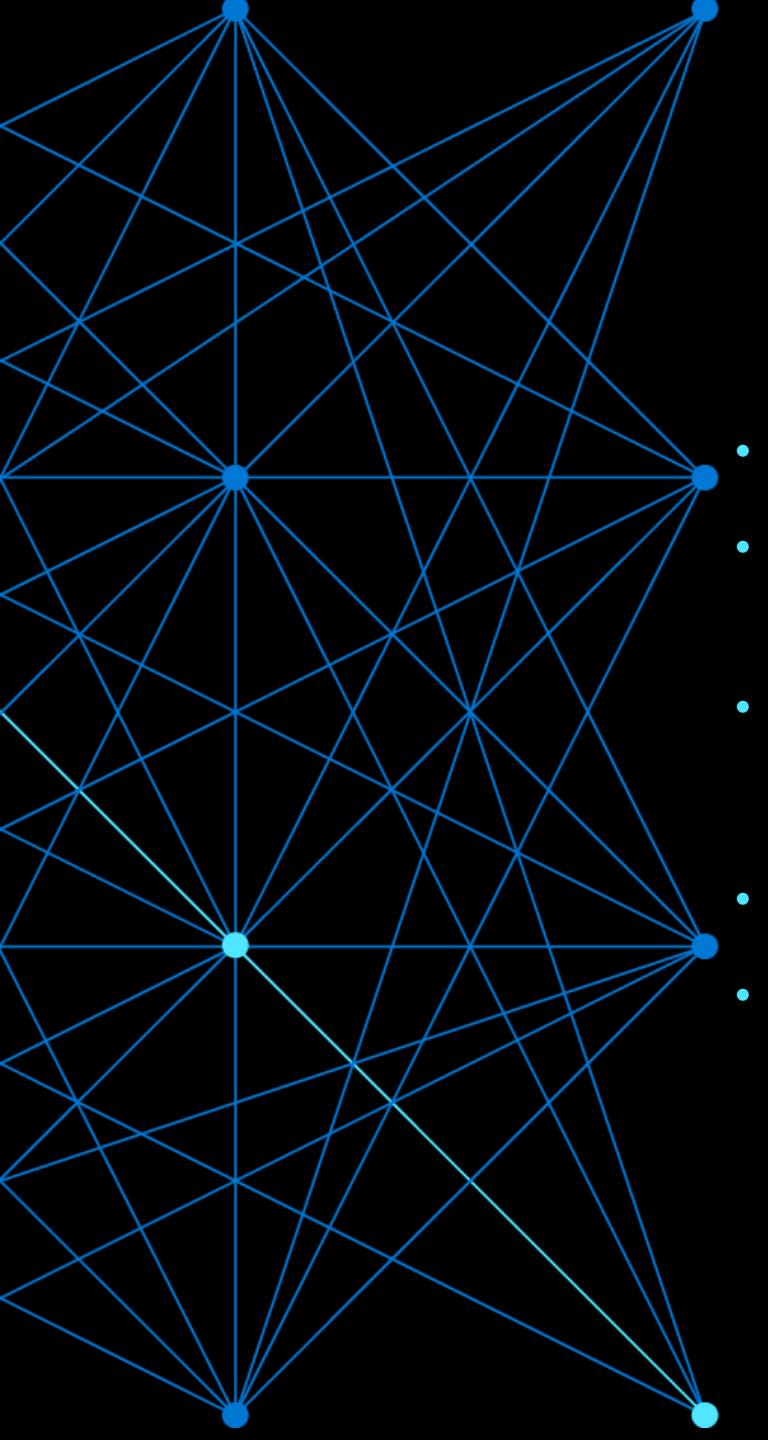




Azure Redis Cache



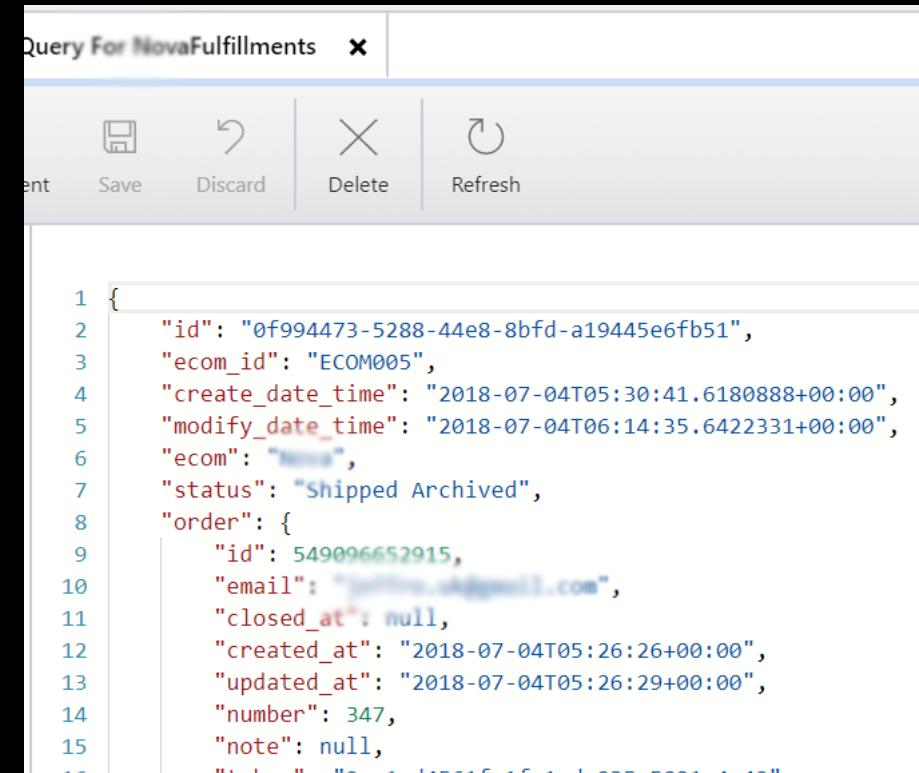
- Platform-as-a-Service
- In-memory database
- Ultra-fast performance
- Widely used for content caching



Azure Cosmos DB

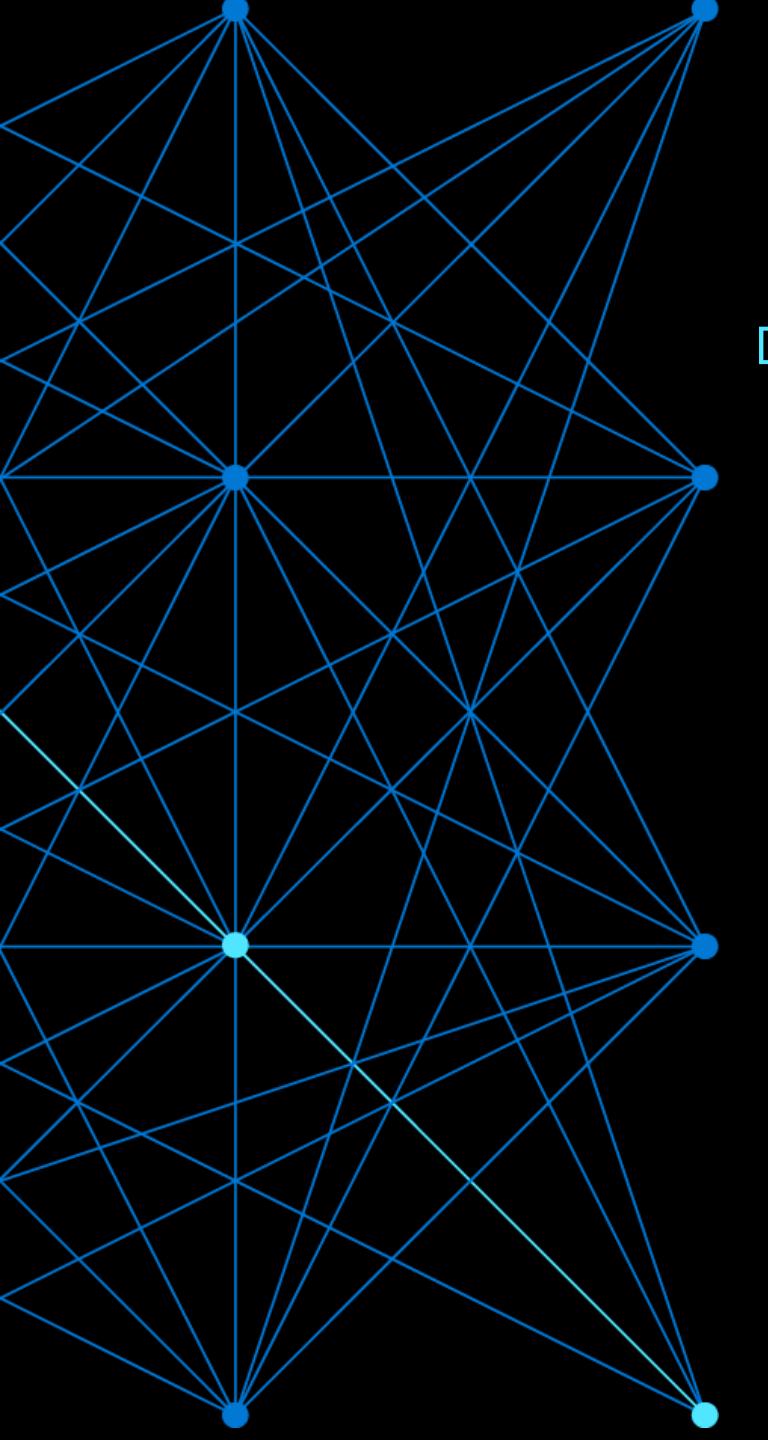


- Platform-as-a-Service
- NoSQL / Document (JSON) database
- Queryable with SQL-like syntax
 - Mongo, Cassandra APIs
- Globally distributed
- Low latency for transactional workloads



```
Query For NovaFulfillments x
Save Discard Delete Refresh

1 {
2   "id": "0f994473-5288-44e8-8bfd-a19445e6fb51",
3   "ecom_id": "ECOM005",
4   "create_date_time": "2018-07-04T05:30:41.6180888+00:00",
5   "modify_date_time": "2018-07-04T06:14:35.6422331+00:00",
6   "ecom": "████████",
7   "status": "Shipped Archived",
8   "order": {
9     "id": 549006652915,
10    "email": "████████@████.com",
11    "closed_at": null,
12    "created_at": "2018-07-04T05:26:26+00:00",
13    "updated_at": "2018-07-04T05:26:29+00:00",
14    "number": 347,
15    "note": null,
16    "url": "https://www.████.com/track/549006652915"
17  }
18}
```



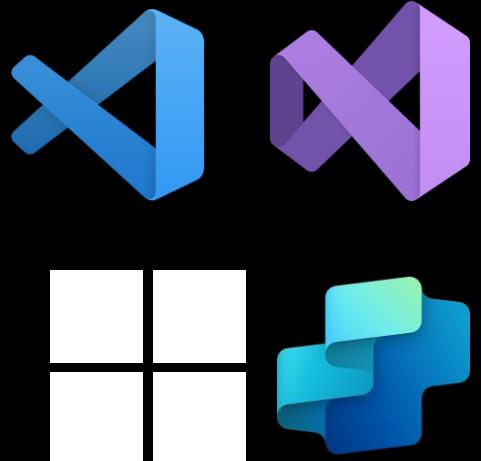
Challenge 02

Discussion Points

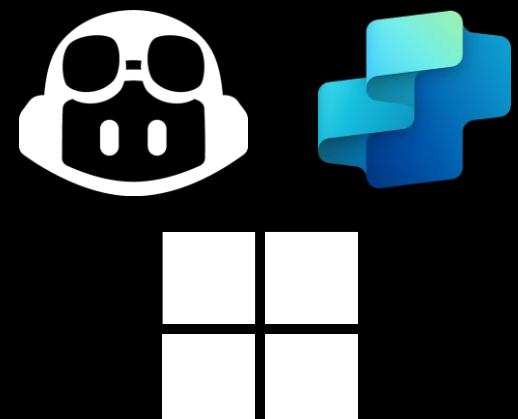
- GitHub Copilot
 - Agent mode
 - Model selection
 - Context window
- Prompts
- MCP
 - Host, Client, Server

MCP is an open protocol that standardizes how applications provide context to LLMs.

Hosts



Clients



Servers



MCP Deep-Dive

Adapted from Mahesh Murag (Anthropic) AI Engineer session

MCP Client

Invokes Tools

Queries for Resources

Interpolates Prompts

MCP Server

Exposes Tools

Exposes Resources

Exposes Prompts

Tools

Model-controlled

Functions invoked by
the model

Retrieve / search

Send a message

Update DB records

Resources

Application-controlled

Data exposed to the
application

Files

Database records

API Responses

Prompts

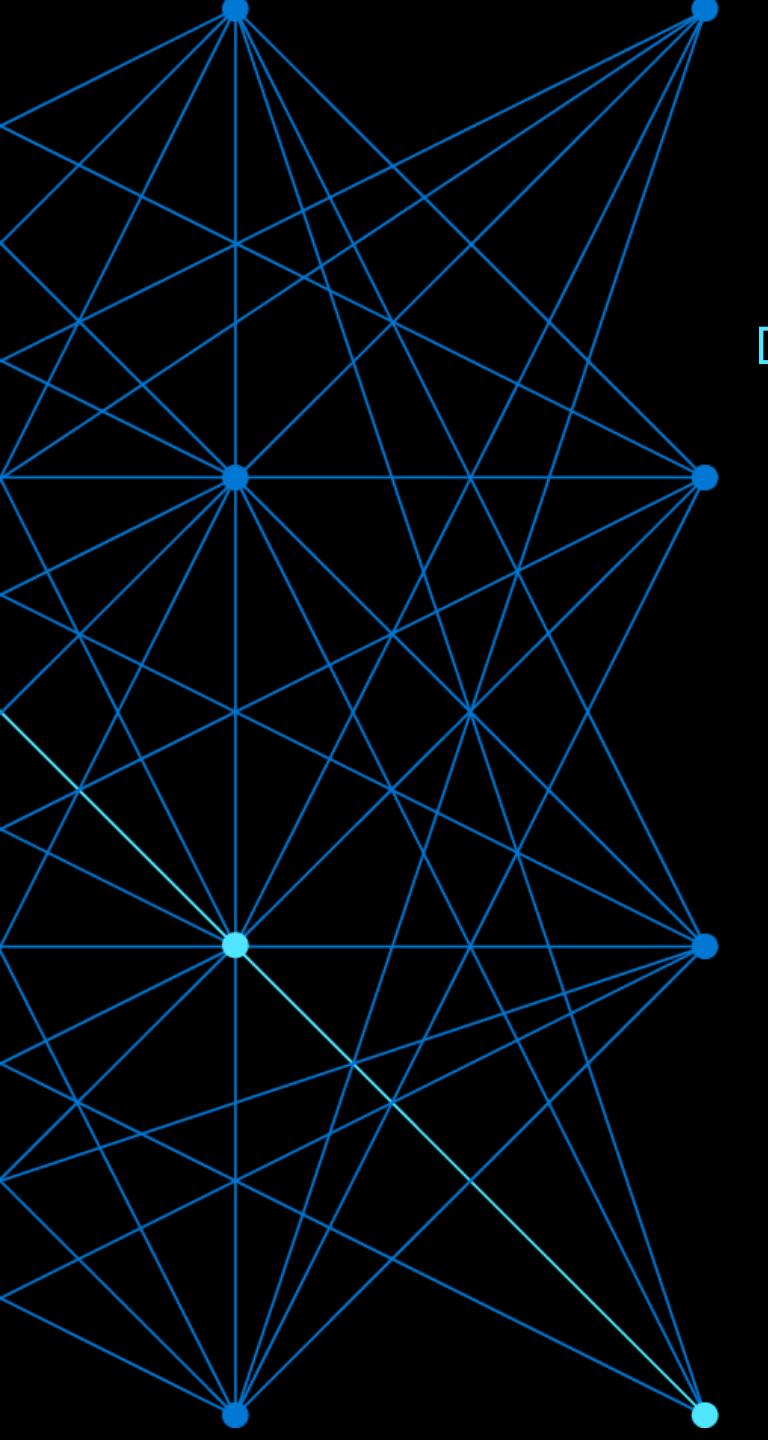
User-controlled

Pre-defined templates
for AI interactions

Document Q&A

Transcript Summary

Output as JSON



Challenge 03

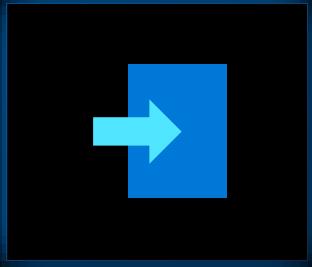
Discussion Points

- Document Intelligence
 - **Tip:** Extractor models take 10+ minutes each to train
- Service Bus and Messaging

How does document process automation work?

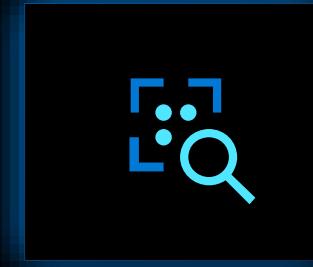
Azure Cognitive Services can help ingest, extract and enrich data

INGEST



Utilize Document Intelligence for digitizing printed or handwritten documents

EXTRACT



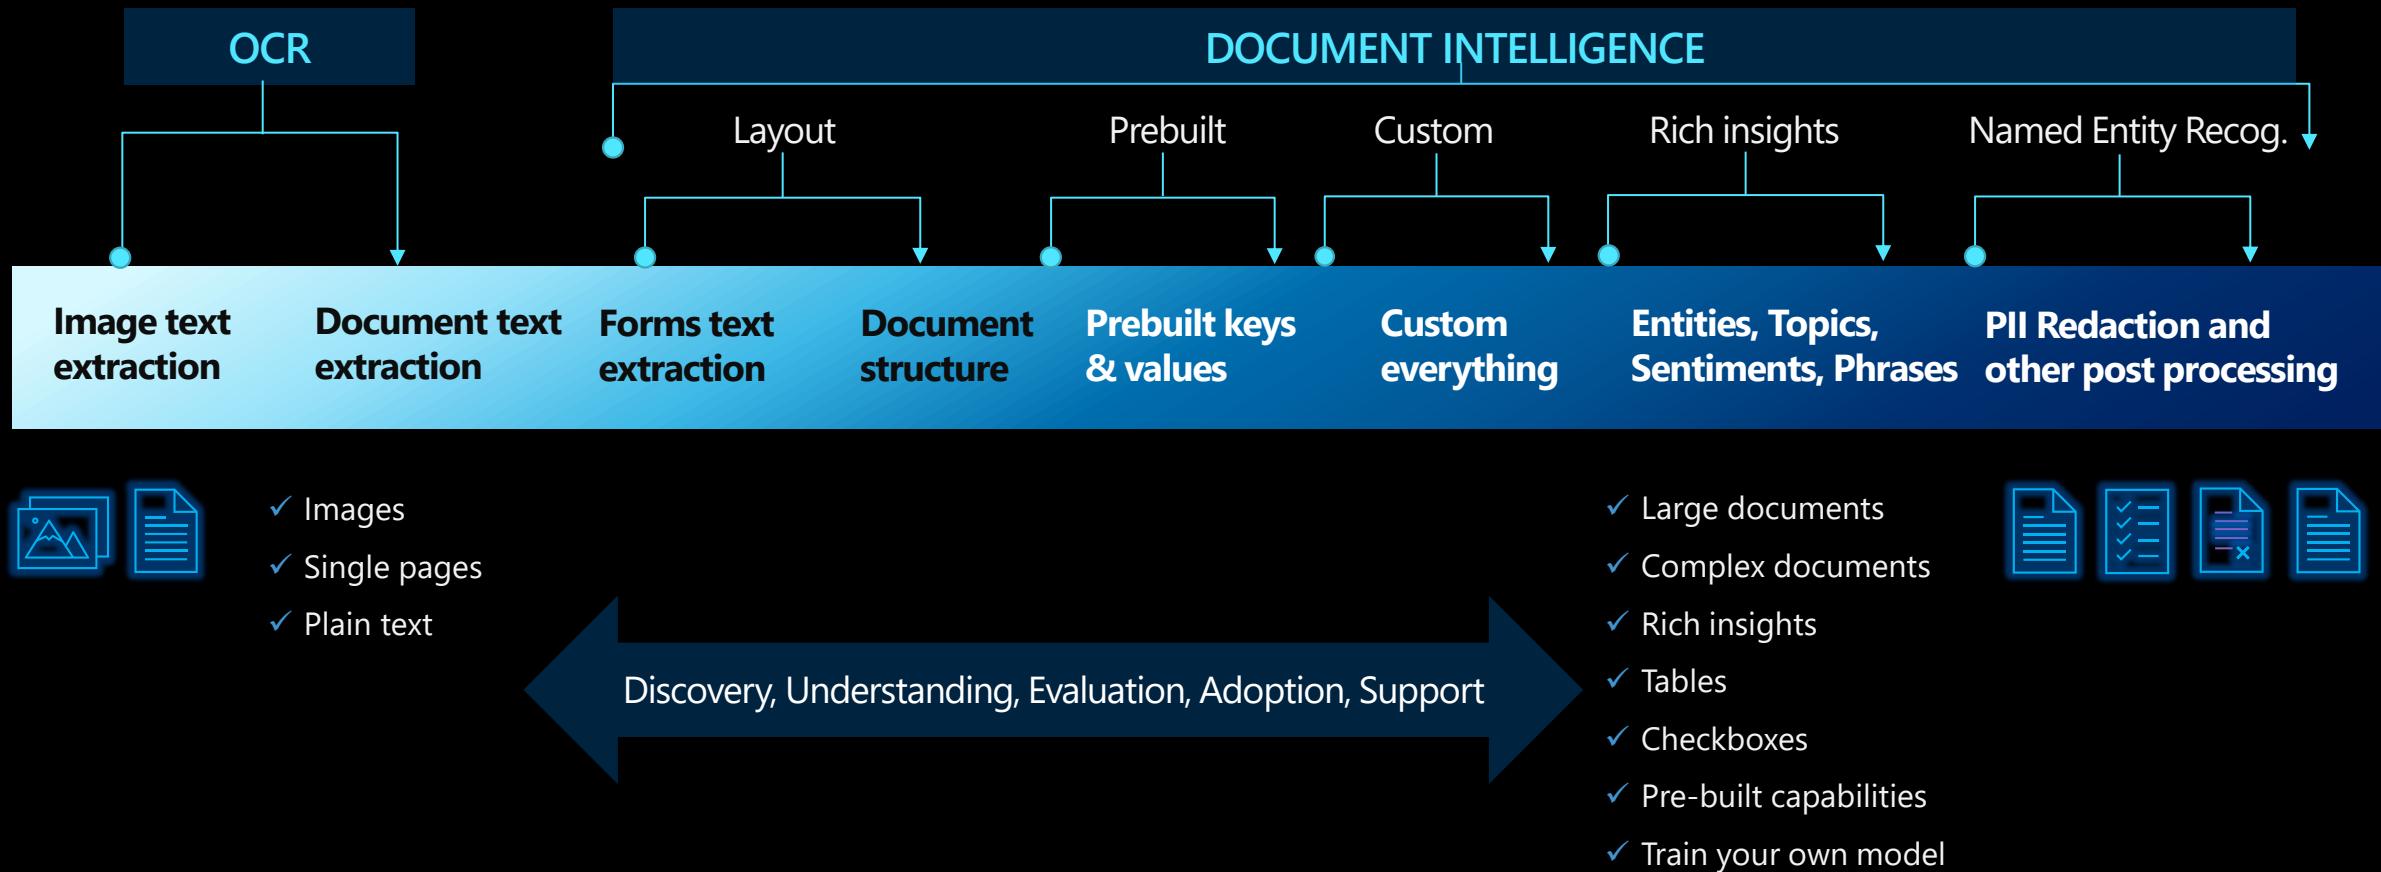
Use pre-trained or custom AI models for entity and key pair extraction with Document Intelligence

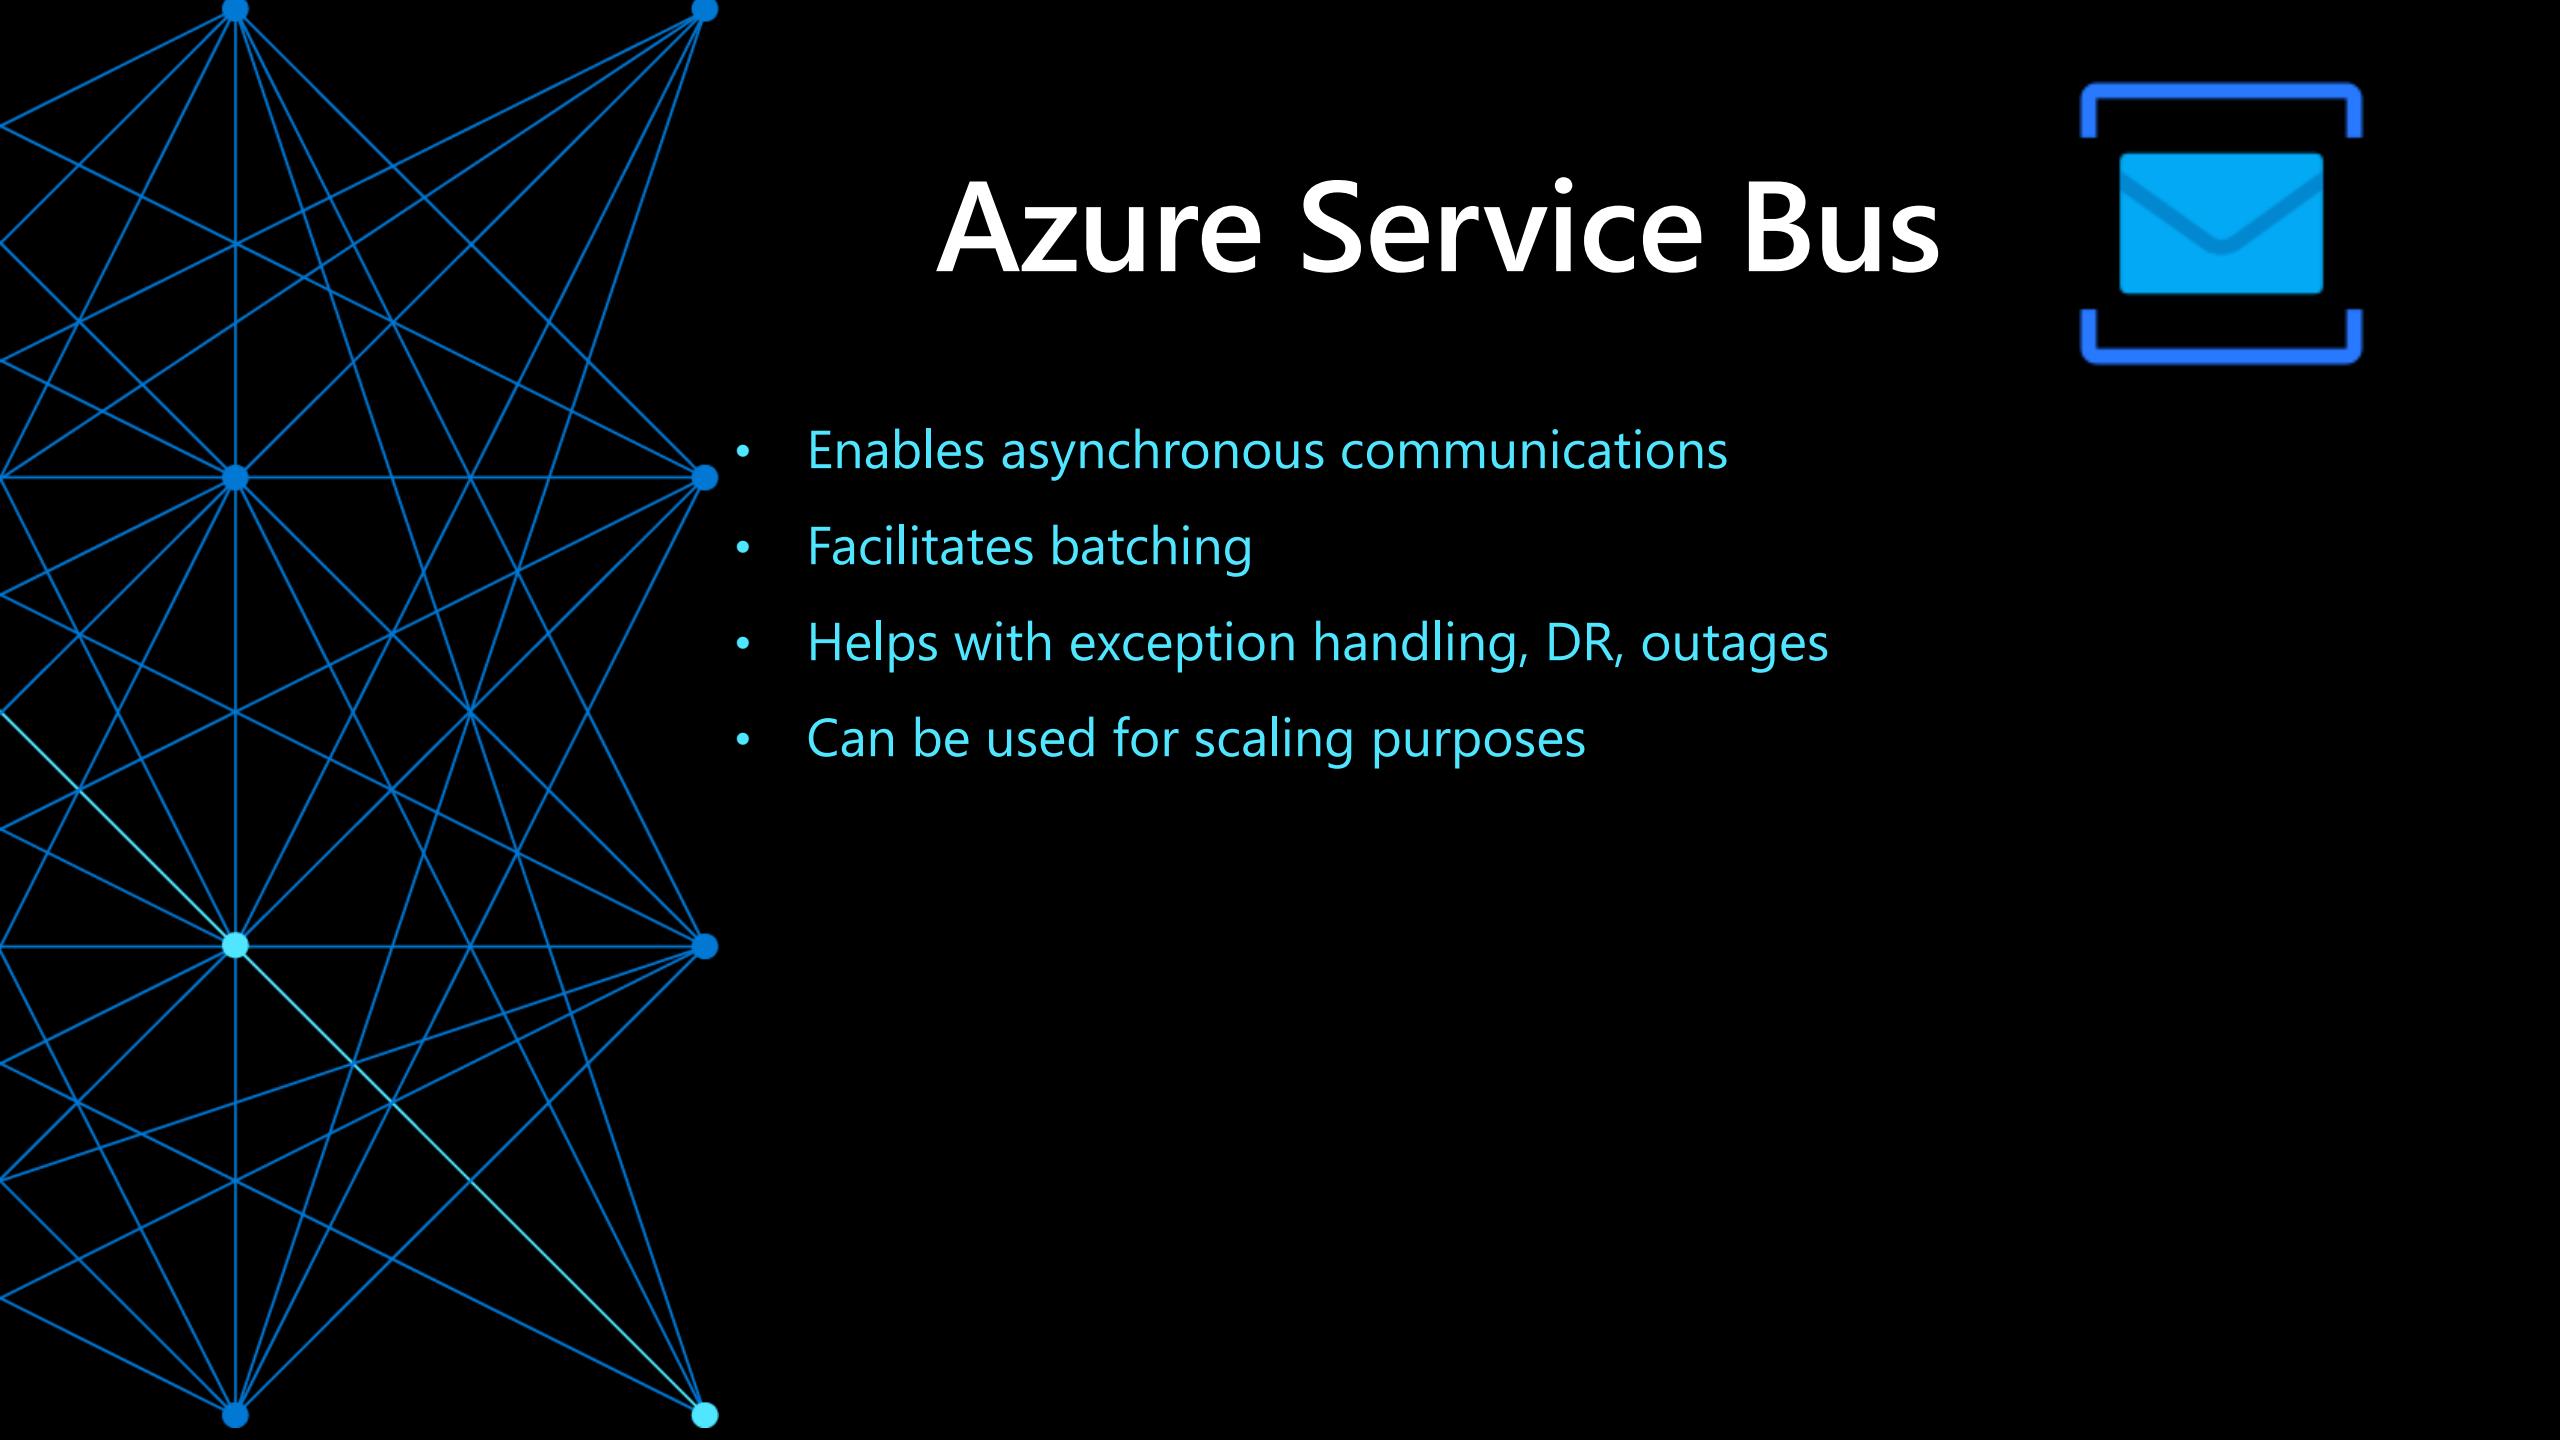
ENRICH



Use prebuilt AI services like Text Analytics or custom LUIS for documents models to enrich the data and make it relevant to your business needs

OCR and Document Intelligence

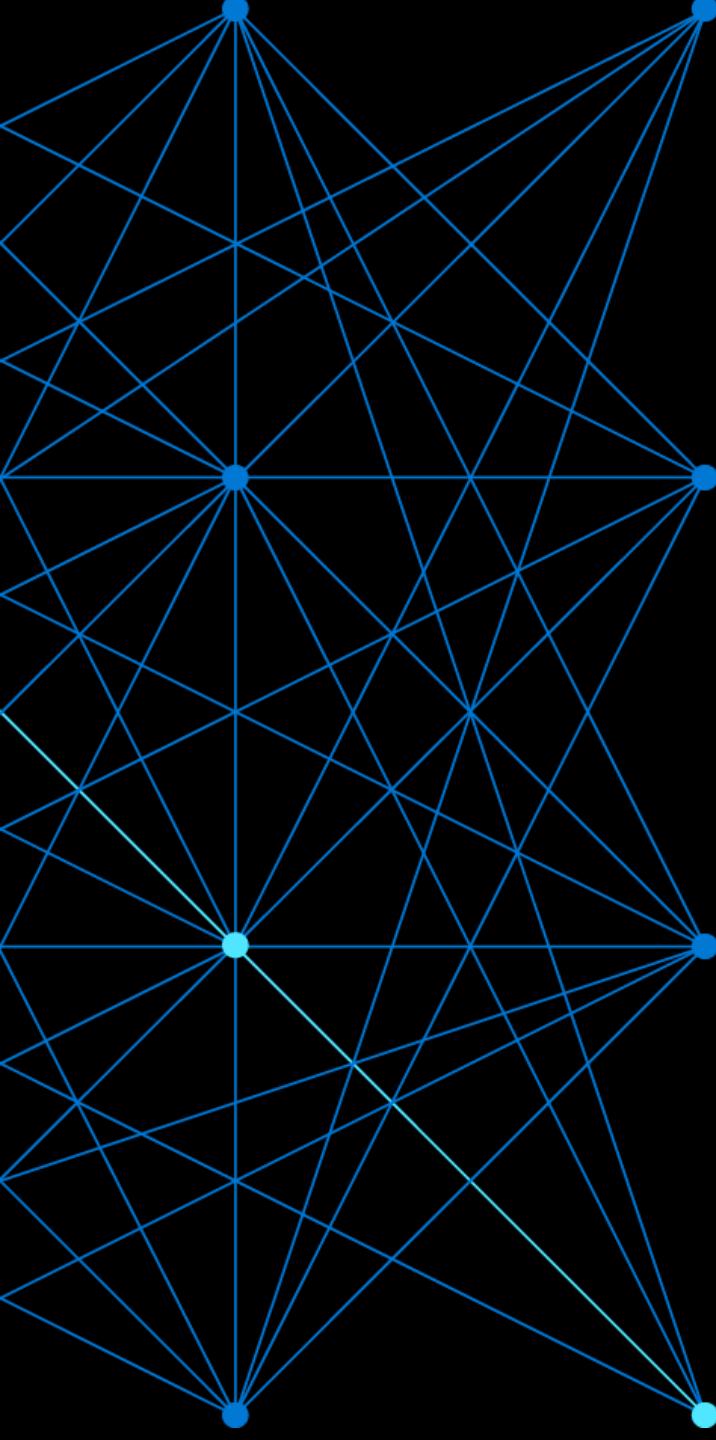




Azure Service Bus



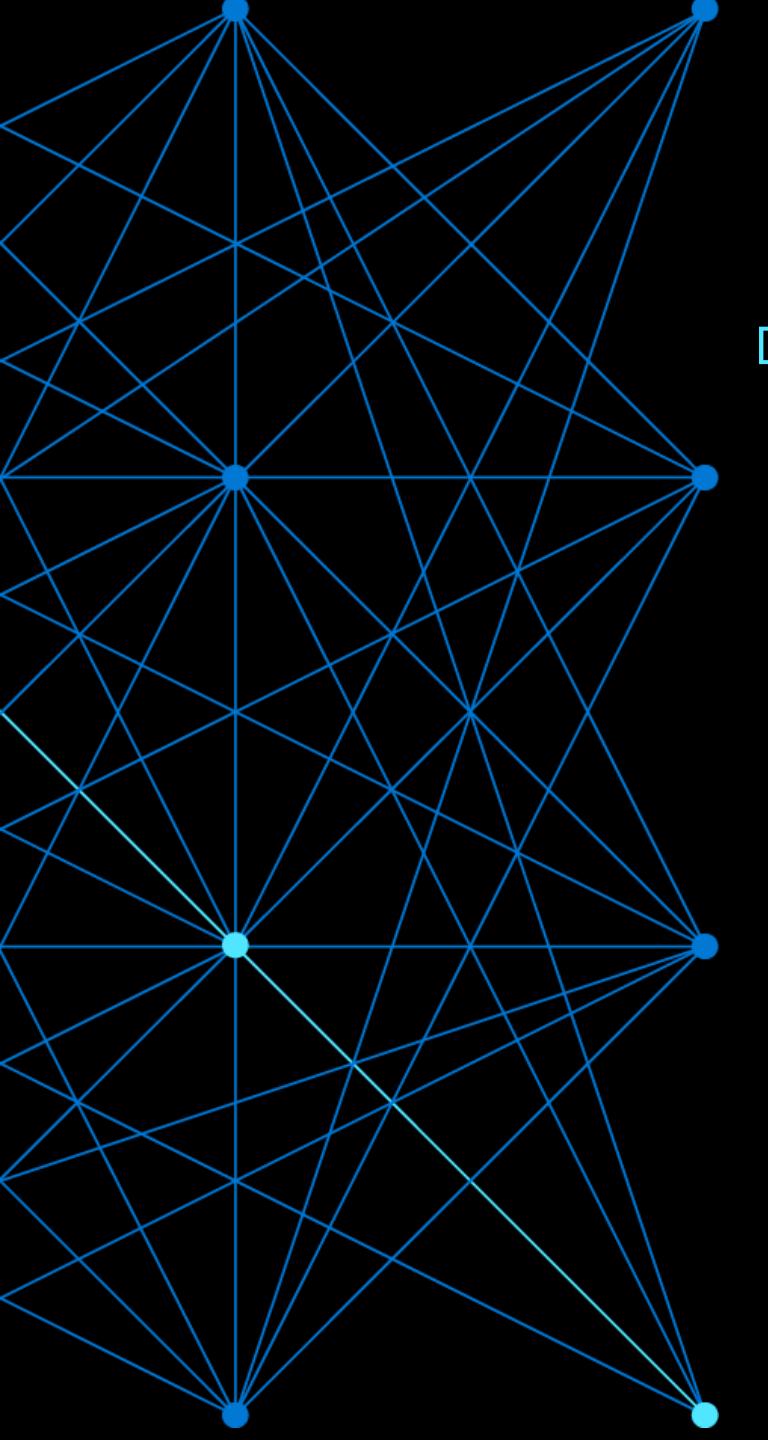
- Enables asynchronous communications
- Facilitates batching
- Helps with exception handling, DR, outages
- Can be used for scaling purposes



Challenge 04

Discussion Points

- Quota management
 - Via Azure OpenAI and external mechanisms



Challenge 05

Discussion Points

- Caching
- OpenAI pricing
- App Insights and distributed application monitoring